

Neural Machine Translation

CSC401/2511 A2 Tutorial 1

Outline

1. Typos in the starter code
2. Walkthrough of the assignment
 - a. Overview
 - b. Calculating BLEU scores
 - c. Encoder
 - d. DecoderWithoutAttention
 - e. DecoderWithAttention
 - f. DecoderWithMultiHeadAttention
 - g. Putting it together: EncoderDecoder
 - h. Training and testing loop
3. teach.cs with GPU: srun
4. Q&A

Most of the material is covered in the [SMT lecture](#)

Page number to the slides: [pXX]

Typos in the Starter Code (piazza @354)

1. `EncoderDecoder.init_modules(...):`

`# ...`

`# 3. You will need the following object attributes:`

`# ... self.heads`

2. `BLEU_score(reference, hypothesis, n):`



`BLEU_score(reference, candidate, n):`

3. `DecoderWithAttention.attend(...):`

`c_t : torch.FloatTensor`

A float tensor of shape ``(M, **self.hidden_state_size**)``.
The context vector `c_t` is the product of weights `alpha_t` and `h`.

Overview: Canadian Hansards



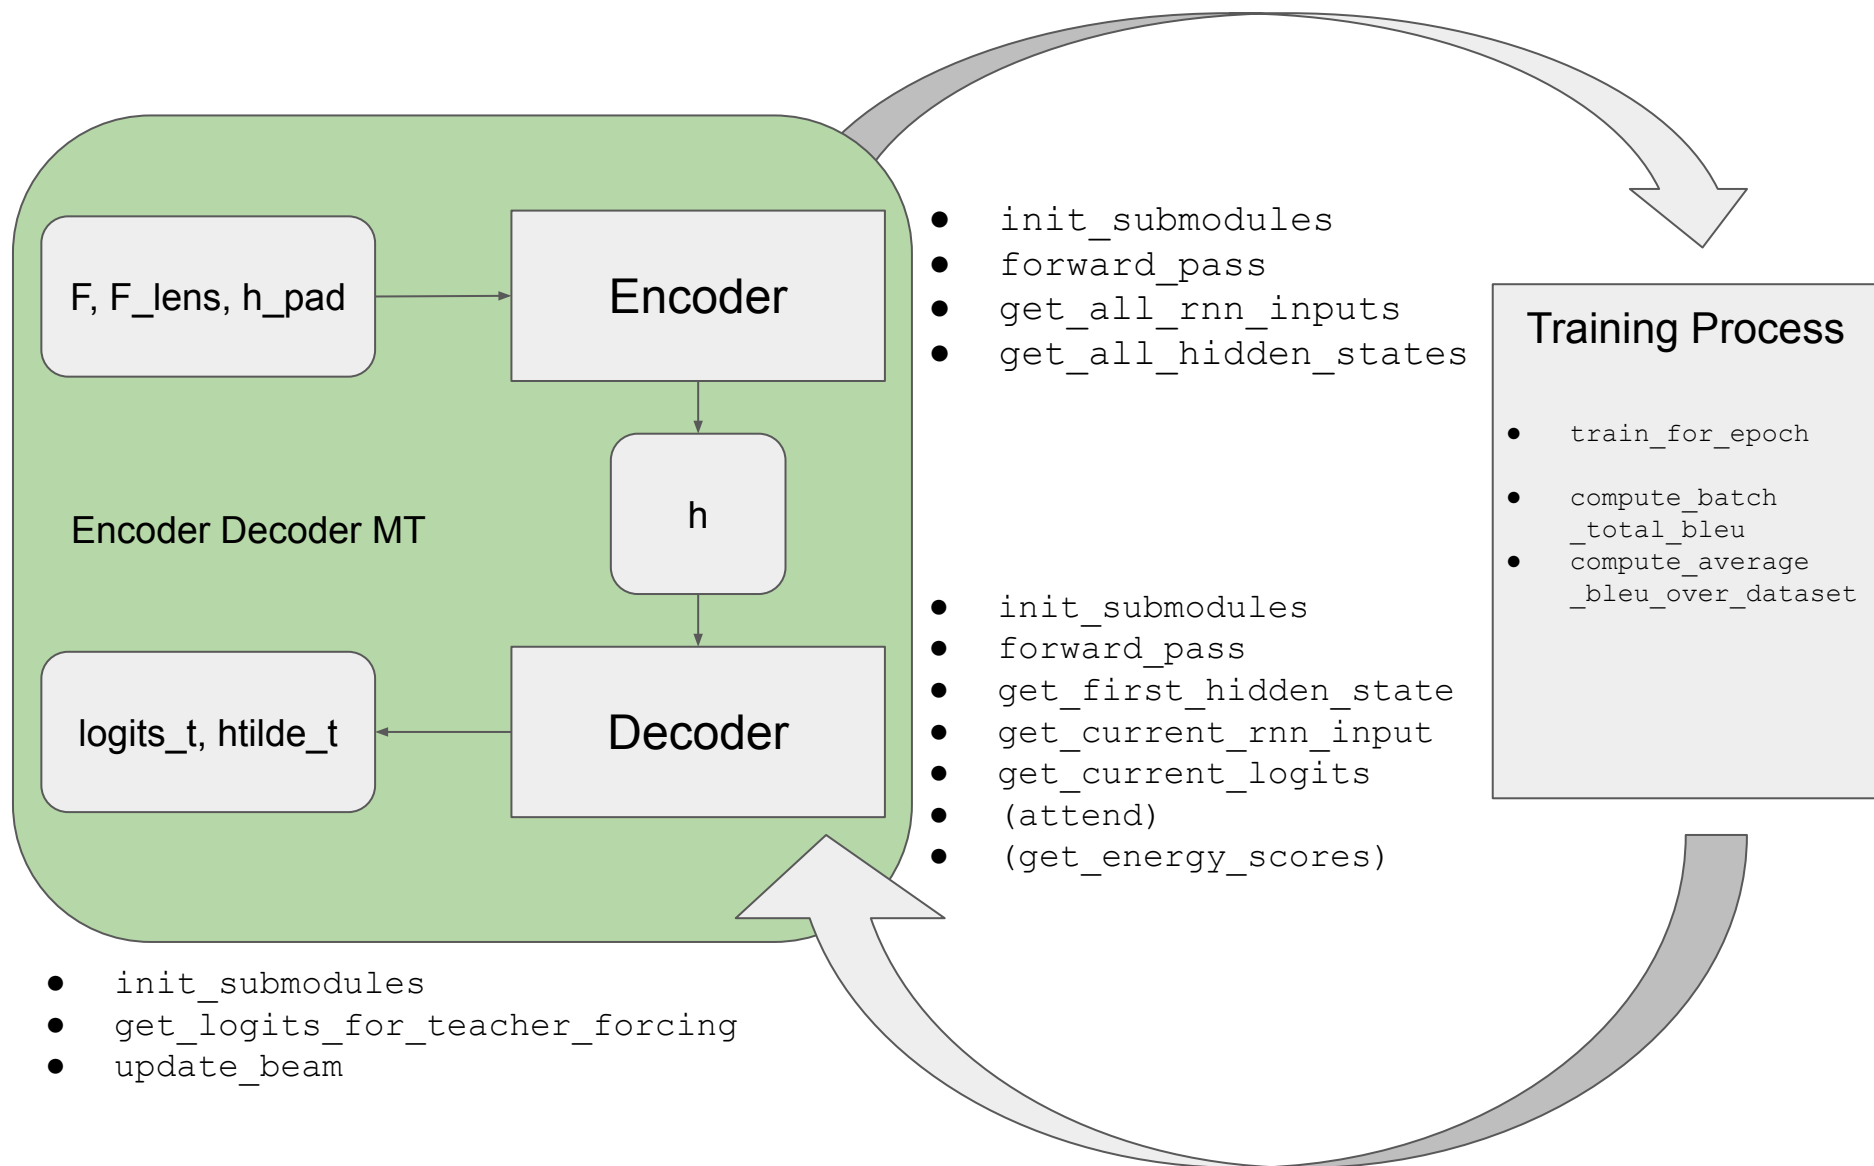
The hon. member is no idiot.
He is a very clever man.

Le depute n'est pas un idiot.
Il est tres intelligent.

- Data located at `/u/cs401/A2/data/Hansard/`
- You don't need to download the data, just set these environment variables

```
export TRAIN=/h/u1/cs401/A2/data/Hansard/Training/  
export TEST=/h/u1/cs401/A2/data/Hansard/Testing/
```

Overview: The Assignment



Calculating BLEU scores

- grouper

```
>>> grouper(['a', 'b', 'c'], 2)
[['a', 'b'], ['b', 'c']]
```

- n_gram_precision

- $\frac{C}{N}$
- [p68]

- brevity_penalty

- $\text{Brevity}_i = \frac{r_i}{c_i}$

$$\text{BP} = \begin{cases} 1 & \text{brevity}_i < 1 \\ e^{1-\text{brevity}_i} & \text{brevity}_i \geq 1 \end{cases}$$

- [p72]

- BLEU_score

- $\overline{BLEU}_c = BP_c \times (p_1 p_2 \dots p_n)^{\frac{1}{n}}$
- [p73]

- No capping
- Only 1 reference and 1 candidate
- Don't include SOS and EOS

Encoder

Encoder

$$x_s = T_F(F_s)$$

$$h = f(x)$$

h: last hidden state of RNN
[p39]

DecoderWithoutAttention

Decoder

$$\tilde{x}_t = T_E(E_{t-1})$$

$$\tilde{h}_t = g(\tilde{x}_t, \tilde{h}_{t-1})$$

$$\text{logits}_t = f(\tilde{h}_t)$$

DecoderWithoutAttention

Decoder

$$\tilde{x}_t = T_E(E_{t-1})$$

$$\tilde{h}_t = g(\tilde{x}_t, \tilde{h}_{t-1})$$

$$\text{logits}_t = f(\tilde{h}_t)$$

- `logits_t` is the **un-normalized log probability**, which means the actual probability should be normalized by **softmax**, instead of average.
- `Pr_b(i) = softmax(logits_t[m])`
- [p40]

DecoderWithAttention

Decoder

$$\tilde{x}_t = [T_E(E_{t-1}), c_{t-1}]$$

$$c_{t-1} = \text{Attend}(h_{t-1}, h_{1:s})$$

$$= \sum_s \alpha_{t-1,s} h_s$$

$$\alpha_{t-1} = \text{softmax}(e_{t-1,1:s}, s)$$

$$e_{t-1,s} = \text{cosine_similarity}(\tilde{h}_{t-1}, h_s)$$

DecoderWithAttention

Decoder

$$\begin{aligned}\tilde{x}_t &= [T_E(E_{t-1}), c_{t-1}] \\ c_{t-1} &= \text{Attend}(h_{t-1}, h_{1:s}) \\ &= \sum_s \alpha_{t-1,s} h_s \\ \alpha_{t-1} &= \text{softmax}(e_{t-1,1:s}, s) \\ e_{t-1,s} &= \text{cosine_similarity}(\tilde{h}_{t-1}, h_s)\end{aligned}$$

- The function to compute attention weights is provided to you
`get_attention_weights`
- Be careful with **t** vs. **t-1**!
- For the assignment, we'll use cosine similarity as the score function (1.4)
- [p43-45]

DecoderWithMultiHeadAttention

Decoder

$$\tilde{h}_{t-1}^{(n)} = \tilde{W}^{(n)} \tilde{h}_{t-1}$$

$$h_s^{(n)} = W^{(n)} h_s$$

$$c_{t-1}^{(n)} = \text{Attention}(\tilde{h}_{t-1}^{(n)}, h_{1:s}^{(n)})$$

$$\tilde{x}_t = [T_F(E_{t-1}), Qc_{t-1}^{(1:N)}]$$


DecoderWithMultiHeadAttention

Decoder

$$\begin{aligned}\tilde{h}_{t-1}^{(n)} &= \tilde{W}^{(n)} \tilde{h}_{t-1} \\ h_s^{(n)} &= W^{(n)} h_s \\ c_{t-1}^{(n)} &= \text{Attention}(\tilde{h}_{t-1}^{(n)}, h_{1:s}^{(n)}) \\ \tilde{x}_t &= [T_F(E_{t-1}), Qc_{t-1}^{(1:N)}]\end{aligned}$$

- You don't really need to slice the hidden weights
- Try starting with heads=1
- You also don't need for loops!
- [p46]

Putting it together: EncoderDecoder



- `init_submodules`
- `get_logits_for_teacher_forcing`
 - Basic idea: replace `y` with `E`
 - [p42]
- `update_beam`
 - One step of the beam search
 - A greedy update function is provided to you, you can test the rest of the assignment by using the `--greedy` option
 - More beam search on next week's tutorial by Zhewei
 - [p55-60]

Training and Testing Loop

- `train_for_epoch`
 - Follow the instructions in the docstring
 - Don't forget to normalize loss!
 - `tqdm`: easy progress bar
- `compute_batch_total_bleu`
 - `a2_bleu_score.BLEU_score` for a batch of sentences
- `compute_average_bleu_over_dataset`
 - Calculate the average BLEU score of the given dataset
 - **Use** `compute_batch_total_bleu`

teach.cs with GPU: srun

- First make sure your code works in cpu mode! Debugging in CUDA mode is much more difficult
- **Basic usage:** `srun -p csc401 your_regular_command`
 - `srun -p csc2511` if you enrolled in CSC 2511
- Check current queue: `squeue -p csc401`
- Keep training after disconnecting: Use `screen` (A.3)